

semaphore工作原理及其使用案例

一 semaphore工作原理	1
1.1 主要结构体分析	1
1.2 如何获取semaphore	2
1.3 如何释放semaphore	6
二 semaphore使用案例	7

一 semaphore工作原理

我们经常在kernel源码中声明semaphore,如下两种方式:

- 静态申请:

```
● #define DEFINE_SEMAPHORE(name) \
●     struct semaphore name = __SEMAPHORE_INITIALIZER(name, 1)
```

- 动态申请

```
● static inline void sema_init(struct semaphore *sem, int val)
● {
●     static struct lock_class_key __key;
●     *sem = (struct semaphore) __SEMAPHORE_INITIALIZER(*sem, val);
●     lockdep_init_map(&sem->lock.dep_map, "semaphore->lock", &__key, 0);
● }
```

它们唯一的区别就是val的数值不同,静态申请的semaphore val=1,即semaphore只能只有一个线程使用,等待的线程task进入休眠状态.

而动态申请的semaphore,val是用户可以自己设定的.如果val > 1,则可以多个线程同时使用此semaphore,直到val <= 0. 其他使用此semaphore的线程task进入休眠等待状态.

1.1 主要结构体分析

主要涉及到两个结构体,分析如下:

semaphore本身的结构体变量:

```
● /* Please don't access any members of this structure directly */
● struct semaphore {
●     raw_spinlock_t    lock;
●     unsigned int      count;
●     struct list_head   wait_list;
● };
```

使用spin lock来保护成员变量count和wait_list.

semaphore等待者结构体,即需要某个相同semaphore的线程task都会填充下面这个结构体:

```
● /* Functions for the contended case */
●
● struct semaphore_waiter {
●     struct list_head list;
●     struct task_struct *task;
●     bool up;
```

```
• };
```

等待semaphore的线程task全部放入链表中,up参数是在释放semaphore的时候设置为true,下面会分析到.

总结上面两个结构体: semaphore_waiter成员变量list是保存需要获取某个semaphore的线程task信息, 这个某个semaphore是谁,怎么关联起来,就是通过结构体semaphore成员变量wait_list,即只要想获取某个semaphore而没有获取成功,则将此semaphore挂载到list链表头上.在释放semaphore的时候,会从这个链表中获取结构体sema_waiter实体,进而也获取到了task信息了.后面详细分析.

1.2 如何获取semaphore

一般在kernel source code里面都是down来获取semaphore的.有下面几种方式来获取想要的semaphore的:

```
• void down(struct semaphore *sem)
• int down_interruptible(struct semaphore *sem)
• int down_killable(struct semaphore *sem)
• int down_trylock(struct semaphore *sem)
• int down_timeout(struct semaphore *sem, long timeout)
```

分别解释如下:

1. down最简单的一种,如果task获取此semaphore,则把此task置为休眠状态,直到semaphore释放,其实就是设置此task调度状态为:TASK_INTERRUPTIBLE
2. down_interruptible是down的一种变种,区别是获取semaphore失败的task可以进入可中断的休眠状态.设置此task的调度状态为:TASK_INTERRUPTIBLE
3. down_killable也是down的变种,如果没有获取semaphore的task进入休眠状态,可被致命信号中断.
4. down_trylock其实是直接判断结构体semaphore成员变量count的数值是否<0
5. down_timeout是down的另一个变种,是获取semaphore失败的task,进入休眠时间timeout之后还没有获取semaphore,则报error.

下面开始讲解上面的实现原理:

```
1. /**
2.  * down - acquire the semaphore
3.  * @sem: the semaphore to be acquired
4.  *
5.  * Acquires the semaphore. If no more tasks are allowed to acquire the
6.  * semaphore, calling this function will put the task to sleep until the
7.  * semaphore is released.
8.  *
9.  * Use of this function is deprecated, please use down_interruptible() or
10. * down_killable() instead.
11. */
12. void down(struct semaphore *sem)
13. {
14.     unsigned long flags;
15.
16.     raw_spin_lock_irqsave(&sem->lock, flags);
17.     if (likely(sem->count > 0))
18.         sem->count--;
19.     else
20.         __down(sem);
21.     raw_spin_unlock_irqrestore(&sem->lock, flags);
```

```

22. }
23. EXPORT_SYMBOL(down);
24.
25. /**
26.  * down_interruptible - acquire the semaphore unless interrupted
27.  * @sem: the semaphore to be acquired
28.  *
29.  * Attempts to acquire the semaphore. If no more tasks are allowed to
30.  * acquire the semaphore, calling this function will put the task to sleep.
31.  * If the sleep is interrupted by a signal, this function will return
    -EINTR.
32.  * If the semaphore is successfully acquired, this function returns 0.
33.  */
34. int down_interruptible(struct semaphore *sem)
35. {
36.     unsigned long flags;
37.     int result = 0;
38.
39.     raw_spin_lock_irqsave(&sem->lock, flags);
40.     if (likely(sem->count > 0))
41.         sem->count--;
42.     else
43.         result = __down_interruptible(sem);
44.     raw_spin_unlock_irqrestore(&sem->lock, flags);
45.
46.     return result;
47. }
48. EXPORT_SYMBOL(down_interruptible);
49.
50. /**
51.  * down_killable - acquire the semaphore unless killed
52.  * @sem: the semaphore to be acquired
53.  *
54.  * Attempts to acquire the semaphore. If no more tasks are allowed to
55.  * acquire the semaphore, calling this function will put the task to sleep.
56.  * If the sleep is interrupted by a fatal signal, this function will return
57.  * -EINTR. If the semaphore is successfully acquired, this function returns
58.  * 0.
59.  */
60. int down_killable(struct semaphore *sem)
61. {
62.     unsigned long flags;
63.     int result = 0;
64.
65.     raw_spin_lock_irqsave(&sem->lock, flags);
66.     if (likely(sem->count > 0))
67.         sem->count--;
68.     else
69.         result = __down_killable(sem);
70.     raw_spin_unlock_irqrestore(&sem->lock, flags);
71.
72.     return result;
73. }
74. EXPORT_SYMBOL(down_killable);

```

```

75.
76. /**
77.  * down_trylock - try to acquire the semaphore, without waiting
78.  * @sem: the semaphore to be acquired
79.  *
80.  * Try to acquire the semaphore atomically. Returns 0 if the semaphore has
81.  * been acquired successfully or 1 if it cannot be acquired.
82.  *
83.  * NOTE: This return value is inverted from both spin_trylock and
84.  * mutex_trylock! Be careful about this when converting code.
85.  *
86.  * Unlike mutex_trylock, this function can be used from interrupt context,
87.  * and the semaphore can be released by any task or interrupt.
88.  */
89. int down_trylock(struct semaphore *sem)
90. {
91.     unsigned long flags;
92.     int count;
93.
94.     raw_spin_lock_irqsave(&sem->lock, flags);
95.     count = sem->count - 1;
96.     if (likely(count >= 0))
97.         sem->count = count;
98.     raw_spin_unlock_irqrestore(&sem->lock, flags);
99.
100.     return (count < 0);
101. }
102. EXPORT_SYMBOL(down_trylock);
103.
104. /**
105.  * down_timeout - acquire the semaphore within a specified time
106.  * @sem: the semaphore to be acquired
107.  * @timeout: how long to wait before failing
108.  *
109.  * Attempts to acquire the semaphore. If no more tasks are allowed to
110.  * acquire the semaphore, calling this function will put the task to
111.  * sleep.
112.  * If the semaphore is not released within the specified number of
113.  * jiffies,
114.  * this function returns -ETIME. It returns 0 if the semaphore was
115.  * acquired.
116.  */
117. int down_timeout(struct semaphore *sem, long timeout)
118. {
119.     unsigned long flags;
120.     int result = 0;
121.
122.     raw_spin_lock_irqsave(&sem->lock, flags);
123.     if (likely(sem->count > 0))
124.         sem->count--;
125.     else
126.         result = __down_timeout(sem, timeout);
127.     raw_spin_unlock_irqrestore(&sem->lock, flags);
128.
129.     return result;
130. }
131. EXPORT_SYMBOL(down_timeout);

```

```

126.     return result;
127. }
128. EXPORT_SYMBOL(down_timeout);
129.
130. static ninline void __sched __down(struct semaphore *sem)
131. {
132.     __down_common(sem, TASK_UNINTERRUPTIBLE, MAX_SCHEDULE_TIMEOUT);
133. }
134.
135. static ninline int __sched __down_interruptible(struct semaphore *sem)
136. {
137.     return __down_common(sem, TASK_INTERRUPTIBLE, MAX_SCHEDULE_TIMEOUT);
138. }
139.
140. static ninline int __sched __down_killable(struct semaphore *sem)
141. {
142.     return __down_common(sem, TASK_KILLABLE, MAX_SCHEDULE_TIMEOUT);
143. }
144.
145. static ninline int __sched __down_timeout(struct semaphore *sem, long
    timeout)
146. {
147.     return __down_common(sem, TASK_UNINTERRUPTIBLE, timeout);
148. }

```

通过上面的代码获取semaphore的时候,分两个处理逻辑:

1. 获取semaphore成功,直接对count减一操作
2. 获取semaphore失败,根据对获取失败semaphore进程的休眠状态分别处理

有必要非常注意的是,在整个的获取semaphore的过程中都被spin lock 保护起来了.使用raw_spin_lock_irqsave来关闭本地CPU的中断.

我们来看它们的通用核心函数:__down_common:

```

• /* Functions for the contended case */
• /* 保存获取失败semaphore的进程信息 */
• struct semaphore_waiter {
•     struct list_head list;
•     struct task_struct *task;
•     bool up;
• };
•
• /*
•  * Because this function is inlined, the 'state' parameter will be
•  * constant, and thus optimised away by the compiler. Likewise the
•  * 'timeout' parameter for the cases without timeouts.
•  */
• static inline int __sched __down_common(struct semaphore *sem, long state,
    long timeout)
• {
•     /*获取semaphore失败的进程*/
•     struct task_struct *task = current;
•     struct semaphore_waiter waiter;
•     /*填充等待semaphore进程信息结构体变量,失败的全部放在一个链表中*/
•     list_add_tail(&waiter.list, &sem->wait_list);
•     waiter.task = task;
•     /*获取semaphore失败的进程,up变量设置为false,在是否semaphore的时候会置为true*/
•     waiter.up = false;

```

```

• /*根据进程设定的状态信息执行不同的调度操作*/
• for (;;) {
•     if (signal_pending_state(state, task))
•         goto interrupted;
•     if (unlikely(timeout <= 0))
•         goto timed_out;
•     /*设置进程状态为state*/
•     __set_task_state(task, state);
•     /*必须关闭spin lock,因为调用__down_common函数的时候,已经调用了
• raw_spin_lock_irqsave spin lock*/
•     raw_spin_unlock_irq(&sem->lock);
•     /*释放cpu,进入休眠状态.如果不释放spin lock而直接进入休眠状态,系统会panic*/
•     timeout = schedule_timeout(timeout);
•     raw_spin_lock_irq(&sem->lock);
•     if (waiter.up)
•         return 0;
• }
• /*timeout之后还没有获取semaphore,将此task waiter信息清空*/
• timed_out:
•     list_del(&waiter.list);
•     return -ETIME;
• /*休眠被中断唤醒,将此task的waiter信息从链表中清空*/
• interrupted:
•     list_del(&waiter.list);
•     return -EINTR;
• }

```

1.3 如何释放semaphore

我们经常看到释放semaphore使用up函数,那么它的原理是什么呢?

```

• /**
•  * up - release the semaphore
•  * @sem: the semaphore to release
•  *
•  * Release the semaphore. Unlike mutexes, up() may be called from any
•  * context and even by tasks which have never called down().
•  */
• void up(struct semaphore *sem)
• {
•     unsigned long flags;
•
•     raw_spin_lock_irqsave(&sem->lock, flags);
•     /*如果waiter 链表为空,则将count++,即大部分情况下,count数值为1*/
•     if (likely(list_empty(&sem->wait_list)))
•         sem->count++;
•     else
•         __up(sem); /*获取waiter 链表成员,并唤醒此链表成员上的进程*/
•     raw_spin_unlock_irqrestore(&sem->lock, flags);
• }
• EXPORT_SYMBOL(up);
•
• static noinline void __sched __up(struct semaphore *sem)

```

```

• { /*获取链表上第一个节点*/
•     struct semaphore_waiter *waiter = list_first_entry(&sem->wait_list,
•                                     struct semaphore_waiter, list);
•     /*将执行释放semaphore的链表元素从waiter链表中清除*/
•     list_del(&waiter->list);
•     waiter->up = true; /*设置flag为true*/
•     wake_up_process(waiter->task); /*直接唤醒此进程,参与调度*/
• }

```

从up函数的解释很有意思:

不像mutex, 释放semaphore,可以在任意的上下文,甚至没有调用down函数的进程都可以主动去释放semaphore.

整体比较简单

二 semaphore使用案例

下面写一个使用semaphore的例子如下:

```

1. #include <linux/init.h>
2. #include <linux/module.h>
3. #include <linux/stat.h>
4. #include <linux/kdev_t.h>
5. #include <linux/fs.h>
6. #include <linux/device.h>
7. #include <linux/cdev.h>
8. #include <asm/uaccess.h>
9. #include <linux/delay.h>
10. #include <linux/semaphore.h>
11.
12. static dev_t devid;
13. static struct class *cls = NULL;
14. static struct cdev mydev;
15. /*静态和动态申请,都可以*/
16. //static DEFINE_SEMAPHORE(mysema);
17. static struct semaphore mysema;
18.
19. static int my_open(struct inode *inode, struct file *file)
20. {
21.     int i;
22.     /*获取semaphore成功则,每隔1s打印一次log*/
23.     while(down_interruptible(&mysema) != 0);
24.     for(i = 0; i < 10; i++) {
25.         printk("semaphore test:%d\n", i);
26.         ssleep(1);
27.     }
28.     /*释放semaphore.后面每次open节点都会再次获取semaphore mysema*/
29.     up(&mysema);
30.     printk("open success!\n");
31.     return 0;
32. }
33.
34. static int my_release(struct inode *inode, struct file *file)
35. {

```

```

36.     printk("close success!\n");
37.     return 0;
38. }
39.
40. //定义文件操作
41. static struct file_operations myfops = {
42.     .owner = THIS_MODULE,
43.     .open = my_open,
44.     .release = my_release,
45. };
46.
47. static void hello_cleanup(void)
48. {
49.     cdev_del(&mydev);
50.     device_destroy(cls, devid);
51.     class_destroy(cls);
52.     unregister_chrdev_region(devid, 1);
53. }
54.
55. static __init int hello_init(void)
56. {
57.     int result;
58.     /*动态创建一个semaphore mysema,并且count设置为1*/
59.     sema_init(&mysema, 1);
60.
61.     //动态注册设备号
62.     if(( result = alloc_chrdev_region(&devid, 0, 1, "samar-alloc-dev") ) !=
    0) {
63.         printk("register dev id error:%d\n", result);
64.         goto err;
65.     } else {
66.         printk("register dev id success!\n");
67.     }
68.     //动态创建设备节点
69.     cls = class_create(THIS_MODULE, "samar-class");
70.     if(IS_ERR(cls)) {
71.         printk("create class error!\n");
72.         goto err;
73.     }
74.
75.     if(device_create(cls, NULL, devid, "", "hello%d", 0) == NULL) {
76.         printk("create device error!\n");
77.         goto err;
78.     }
79.     //字符设备注册
80.     mydev.owner = THIS_MODULE;           //必要的成员初始化
81.     mydev.ops = &myfops;
82.     cdev_init(&mydev, &myfops);
83.     //添加一个设备
84.     result = cdev_add(&mydev, devid, 1);
85.     if(result != 0) {
86.         printk("add cdev error!\n");
87.         goto err;
88.     }

```



```

89.
90.     printk(KERN_ALERT "hello init success!\n");
91.     return 0;
92.err:
93.     hello_cleanup();
94.     return -1;
95.}
96.
97.static __exit void hello_exit(void)
98.{
99.     hello_cleanup();
100.     printk("helloworld exit!\n");
101.}
102.
103. module_init(hello_init);
104. module_exit(hello_exit);
105.
106. MODULE_LICENSE("GPL");
107. MODULE_AUTHOR("samarxie");
108. MODULE_DESCRIPTION("For semaphore use method");

```

semaphore原理比较容易.关键还是使用了spin lock
通过用户空间操作/dev/hello0节点:

```

• #include <stdio.h>
• #include <stdlib.h>
• #include <unistd.h>
• #include <string.h>
•
• #include <sys/types.h>
• #include <sys/stat.h>
• #include <fcntl.h>
•
• int
• main(void)
• {
•
•     int fd;
•     char buf[100];
•     int size;
•
•
•     fd = open("/dev/hello0", O_RDWR);
•     if(!fd) {
•         perror("open");
•         exit(-1);
•     }
•     printf("open success!\n");
•
•     close(fd);
•     return 0;
• }

```


